

Developing high-performance Coroutines for ARMs

Klemens Morgenstern

10.03.2018

Introduction

- Electrical engineer by training
- C++ developer by passion
- boost contributor
- Independent contractor
- Open for consulting

Motivation

- Coroutines are awesome

Motivation

- Coroutines are awesome
- Low overhead alternative to threads

Motivation

- Coroutines are awesome
- Low overhead alternative to threads
- Do not require an OS / Scheduler

Motivation

- Coroutines are awesome
- Low overhead alternative to threads
- Do not require an OS / Scheduler
- More readable alternative to state-machines

Motivation

- Coroutines are awesome
- Low overhead alternative to threads
- Do not require an OS / Scheduler
- More readable alternative to state-machines
- Not well enough known

1. Introduction to coroutines

1. Introduction to coroutines
2. Implementation

1. Introduction to coroutines
2. Implementation
3. Performance comparisons

Definition

- Function with it's own stack

Definition

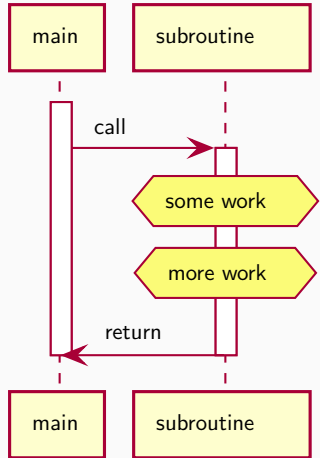
- Function with it's own stack
- can interrupt it's own execution

Definition

- Function with it's own stack
- can interrupt it's own execution
- and be resumed later

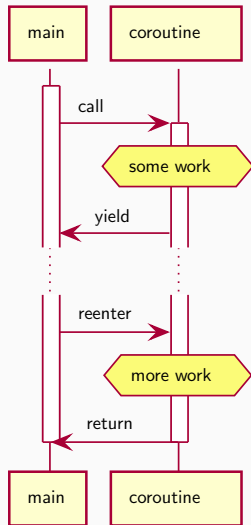
Subroutine

```
1  int main()  
2  {  
3      subroutine();  
4      return 0;  
5  }  
6  
7  void subroutine()  
8  {  
9      some_work();  
10     more_work();  
11 }
```



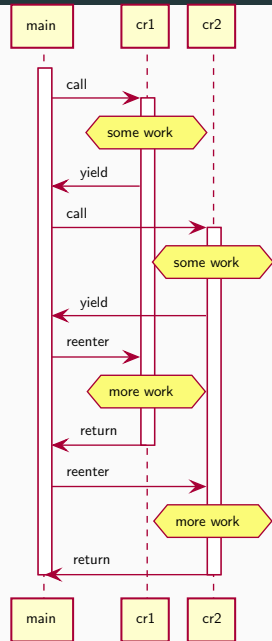
Coroutine

```
1  int main()  
2  {  
3      coroutine cr{&cr_impl};  
4      cr.spawn();  
5      //...main does something  
6      cr.reenter();  
7      return 0;  
8  }  
9  
10 void cr_impl(yield_t &yield)  
11 {  
12     some_work();  
13     yield();  
14     more_work();  
15 }
```



Coroutine concurrency

```
1  int main()  
2  {  
3      coroutine cr1{&cr_impl};  
4      coroutine cr2{&cr_impl};  
5      cr1.spawn();  
6      cr2.spawn();  
7      cr1.reenter();  
8      cr2.reenter();  
9      return 0;  
10 }  
11 void cr_impl(yield_t & yield)  
12 {  
13     some_work();  
14     yield();  
15     more_work();  
16 }
```



Value yield example

```
1  int cr_impl(yield_t & yield)
2  {
3      yield(1);
4      yield(2);
5      return 4;
6  }
7
8  int main()
9  {
10     coroutine<int()> cr{&cr_impl};
11
12     assert(cr.spawn() == 1);
13     assert(cr.reenter() == 2);
14     assert(cr.reenter() == 4);
15     return cr.done() ? 0 : 1;
16 }
```

Argument passing example

```
1  int cr_impl(yield_t & yield , bool)
2  {
3      int value = 1;
4      while(yield(value))
5          value <<= 1;
6      return 0;
7  }
8
9  int main()
10 {
11     coroutine<int(bool)> cr{&cr_impl};
12
13     assert(cr.spawn(true) == 1);
14     assert(cr.reenter(true) == 2);
15     assert(cr.reenter(true) == 4);
16     return cr.reenter(false);
17 }
```

- Deterministic, no real concurrency

Properties

- Deterministic, no real concurrency
- Pass values out on yield in on reenter

Properties

- Deterministic, no real concurrency
- Pass values out on yield in on reenter
- May not be moved in memory

Properties

- Deterministic, no real concurrency
- Pass values out on yield in on reenter
- May not be moved in memory
- Can be stored (e.g. for deep sleep)

Properties

- Deterministic, no real concurrency
- Pass values out on yield in on reenter
- May not be moved in memory
- Can be stored (e.g. for deep sleep)
- Only constraint is the stack-size

Properties

- Deterministic, no real concurrency
- Pass values out on yield in on reenter
- May not be moved in memory
- Can be stored (e.g. for deep sleep)
- Only constraint is the stack-size
- Could be used as callback

Context switch

Main Context

```
1 int main()  
2 {  
3     coroutine cr;  
4     cr.spawn()  
5     /* make_context  
6      *  
7      *  
8      *return from spawn*/;  
9     cr.reenter()  
10    /* switch_context  
11     *  
12     *return from reenter*/;  
13    return 0;  
14 }
```

Coroutine Context

```
1  
2  
3  
4  
5 void cr_impl()  
6 {  
7     yield()  
8     /* switch_context  
9      *  
10    *return from yield*/;  
11 }  
12  
13  
14
```

Context switch

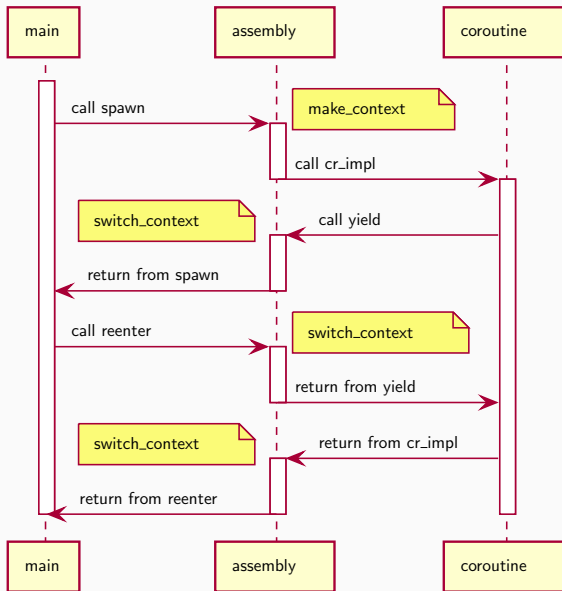
Main Context

```
1 int main()  
2 {  
3     coroutine cr;  
4     cr.spawn()  
5     /* make_context  
6      *  
7      *  
8      *return from spawn*/;  
9     cr.reenter()  
10    /* switch_context  
11     *  
12     *return from reenter*/;  
13    return 0;  
14 }
```

Coroutine Context

```
1  
2  
3  
4  
5 void cr_impl()  
6 {  
7     yield()  
8     /* switch_context  
9      *  
10    *return from yield*/;  
11 }  
12  
13  
14
```

Context switch



Stack Pointer SP

- Points to lowest element on stack

Stack Pointer SP

- Points to lowest element on stack
- Decrement on function entry / incremented on exit (usually)

Stack Pointer SP

- Points to lowest element on stack
- Decrementing on function entry / incrementing on exit (usually)

Stack Pointer SP

- Points to lowest element on stack
- Decrement on function entry / incremented on exit (usually)

C++

```
1 void foo() { }
```

Stack Pointer SP

- Points to lowest element on stack
- Decrement on function entry / incremented on exit (usually)

C++

```
1 void foo() { }
```

ASM (by gcc)

```
1  str fp, [sp, #-4]!  
2  add fp, sp, #0  
3  sub sp, fp, #0  
4  mov a1, a1 @ nop  
5  sub sp, fp, #0  
6  ldr fp, [sp], #4  
7  bx lr
```


Link Register LR

- Points to the code location of function call

Link Register LR

- Points to the code location of function call
- Returning means jump to the location

Link Register LR

- Points to the code location of function call
- Returning means jump to the location
- Pushed on stack for new function call

Link Register LR

- Points to the code location of function call
- Returning means jump to the location
- Pushed on stack for new function call

C++

```
1 void foo() {bar();}
```

Link Register LR

- Points to the code location of function call
- Returning means jump to the location
- Pushed on stack for new function call

C++

```
1 void foo() {bar();}
```

ASM (by gcc)

```
1  push {fp, lr}
2  add fp, sp, #4
3  bl  bar()
4  mov r0, r0 @ nop
5  sub sp, fp, #4
6  pop {fp, lr}
7  bx  lr
```

Argument Register a1-a4

- Used to pass arguments in and return

Argument Register a1-a4

- Used to pass arguments in and return
- Large values store a reference

Argument Register a1-a4

- Used to pass arguments in and return
- Large values store a reference
- a1 stores the return value

Argument Register a1-a4

- Used to pass arguments in and return
- Large values store a reference
- a1 stores the return value
- Only valid in local context, not persistent

Argument Register a1-a4

- Used to pass arguments in and return
- Large values store a reference
- a1 stores the return value
- Only valid in local context, not persistent
- ip is an additional scratch register

Argument Register a1-a4

- Used to pass arguments in and return
- Large values store a reference
- a1 stores the return value
- Only valid in local context, not persistent
- ip is an additional scratch register

C++

```
1 int foo(int x, int y) { return x+y; }
```

Argument Register a1-a4

- Used to pass arguments in and return
- Large values store a reference
- a1 stores the return value
- Only valid in local context, not persistent
- ip is an additional scratch register

C++

```
1 int foo(int x, int y) { return x+y; }
```

ASM (by gcc,-O3)

```
1 add a1, a1, a2  
2 bx lr
```

Variable registers v1-v8

- Load/Store-Architecture

Variable registers v1-v8

- Load/Store-Architecture
- Values must be loaded into variable registers for operations

Variable registers v1-v8

- Load/Store-Architecture
- Values must be loaded into variable registers for operations
- Must be persistent after subroutine calls

Variable registers v1-v8

- Load/Store-Architecture
- Values must be loaded into variable registers for operations
- Must be persistent after subroutine calls
- When used, old values get pushed on the stack

Variable registers v1-v8

- Load/Store-Architecture
- Values must be loaded into variable registers for operations
- Must be persistent after subroutine calls
- When used, old values get pushed on the stack

C++

```
1 void bar(int&, int&,
2         int&, int&);
3 void foo(int &x, int &y,
4         int &z, int &i)
5 {
6     x = i+y+z;
7     bar(x, y, z, i);
8 }
```

Variable registers v1-v8

- Load/Store-Architecture
- Values must be loaded into variable registers for operations
- Must be persistent after subroutine calls
- When used, old values get pushed on the stack

C++

```
1 void bar(int&, int&,  
2         int&, int&);  
3 void foo(int &x, int &y,  
4         int &z, int &i)  
5 {  
6     x = i+y+z;  
7     bar(x,y,z,i);  
8 }
```

ASM (by gcc)

```
1  push {v1, v2, v3, lr}  
2  ldr v2, [a4]  
3  ldr v1, [a2]  
4  ldr lr, [a3]  
5  mla ip, v1, v2, lr  
6  str ip, [a1]  
7  bl bar  
8  pop {v1, v2, v3, lr}  
9  bx lr
```

- `r0 - r3/a1 - a4` Argument/Scratch register

- `r0 - r3 / a1 - a4` Argument/Scratch register
- `r4 - r11 / v1 - v8` Variable registers

- `r0 - r3 / a1 - a4` Argument/Scratch register
- `r4 - r11 / v1 - v8` Variable registers
- `r11 / fp` *Frame pointer on gcc*

- `r0 - r3 / a1 - a4` Argument/Scratch register
- `r4 - r11 / v1 - v8` Variable registers
- `r11 / fp` *Frame pointer on gcc*
- `r12 / ip` Intra procedure scratch

- r0 - r3 / a1 - a4 Argument/Scratch register
- r4 - r11 / v1 - v8 Variable registers
- r11 / fp *Frame pointer on gcc*
- r12 / ip Intra procedure scratch
- r13 / SP Stack pointer

- r0 - r3/a1 - a4 Argument/Scratch register
- r4 - r11/v1 - v8 Variable registers
- r11/ fp *Frame pointer on gcc*
- r12/ ip Intra procedure scratch
- r13/ SP Stack pointer
- r14/ LR Link Register

- r0 - r3 / a1 - a4 Argument/Scratch register
- r4 - r11 / v1 - v8 Variable registers
- r11 / fp *Frame pointer on gcc*
- r12 / ip Intra procedure scratch
- r13 / SP Stack pointer
- r14 / LR Link Register
- r15 / PC Program Counter

Creating a coroutine

```
1 struct coroutine
2 {
3     void * stack_pointer; //valid
4     template<typename Function>
5     Return spawn(Function && func)
6     {
7         return make_context(this, &func, &executor);
8     }
9     static void executor(coroutine * const this_,
10                          Function *func_p)
11     {
12         Function func = *func_p;
13         Return val = func({this_});
14         done = true;
15         switch_context(val, this_);
16     };
17 };
```

Make context

```
1  make_context_0:
2      @_make_context_0(cr* const, void *, void * )
3      @
4      @executor: (cr * const, void *)
5      @
6      @coroutine, func
7      push {v1-v8, lr} @push the link register
8      mov v1, sp      @move the stack pointer to v1
9      ldr sp, [a1]    @set the stack pointer
10     str v1, [a1]    @store the old stack pointer
11
12     bx a3 @call the function
```

Context switch

```
1 void switch_context_0(void*& sp);  
2 void yield (void*& p){switch_context_0(p);}   
3 void reenter(void*& p){switch_context_0(p);}
```

Context switch

```
1 void switch_context_0(void*& sp);
2 void yield (void*& p){switch_context_0(p);}
3 void reenter(void*& p){switch_context_0(p);}
```

```
1 switch_context_0:
2     push {v1-v8, lr} @8
3
4     mov v1, sp @1
5     ldr sp, [a1] @2
6     str v1, [a1] @2
7
8     pop {v1-v8, lr} @8
9     bx lr @1 --> overall 22
```

Context switch uint32_t(uint64_t)

```
1  uint64_t  switch_context_1(uint32_t, void*& sp);
2  uint32_t  switch_context_2(uint64_t, void*& sp);
3  uint64_t  yield  (uint32_t i, void*& p)
4  {
5      return switch_context_1(i, sp);
6  }
7  uint32_t  reenter(uint64_t i, void*& p)
8  {
9      return switch_context_2(i, sp);
10 }
```

Context switch `yield(uint32_t)`

```
1  switch_context_1:
2      @yield(uint32_t, void*) -> uint64_t
3      @      a1      , a2
4      push {v1-v8, lr}
5
6      mov v1, sp
7      ldr sp, [a2]
8      str v1, [a2]
9
10     pop {v1-v8, lr}
11
12     @reenter(uint64_t, void*) -> uint32_t
13     @return value is in a1
14     bx lr
```

Context switch reenter(uint64_t)

```
1  switch_context_2:
2      @reenter(uint64_t, void*) -> uint32_t
3      @          a1-a2      , a3
4      push {v1-v8, lr}
5
6      mov v1, sp
7      ldr sp, [a3]
8      str v1, [a3]
9
10     pop {v1-v8, lr}
11
12     @yield(uint32_t, void*) -> uint64_t
13     @return value is in a1-a2
14     bx lr
```


Statemachine example

```
1 struct statemachine
2 {
3     int state = 0;
4     bool done = false;
5     int operator()()
6     {
7         switch (state)
8         {
9             case 0:
10                 state = 1; return f();
11             case 1:
12                 state = 2; return g();
13             case 2:
14                 done = true; return h();
15         }
16     }
17 };
```

Statemachine example

```
1  int coroutine(embo::yield_t<int()>& yield)
2  {
3      yield(f());
4      yield(g());
5      return h();
6  }
```

Statemachine Assembly

```
1  statemachine::operator()():
2      ldr a4, [a1]    @2
3      push {v1, lr}  @1
4      cmp a4, #0      @1
5      beq .L10        @1
6      cmp a4, #1
7      beq .L3
8      cmp a4, #2
9      beq .L4
10     pop {r4, lr}
11     bx lr
12 .L10:
13     mov a4, #1      @1
14     str a4, [a1]    @2
15     bl f()          @1
16     pop {v1, lr}    @1
17     bx lr           @1 -> 11 for f, 12 for g, 13 for h
```

Coroutine Assembly

```
1  coroutine(yield_t&): @14
2      push {v1, lr} @1
3      mov v1, a1      @1
4      bl f()          @1
5      mov a2, v1      @1
6      bl switch_context_1(int, void*) @1+22; 22 reentry
7      bl g() @1
8      mov a2, v1 @1
9      bl switch_context_1(int, void*) @1+22
10     bl h() @1
11     pop {v1, lr} @1
12     bx lr @1 -> 41 for f, 47 g, 47 for h
```

Statemachine stack example

```
1 struct statemachine {
2     int state = 0;
3     bool use_wchar = false;
4     array<char, 128> buffer;
5     array<wchar_t, 128> wbuffer;
6     void operator>()() {
7         switch (state) {
8             case 0: state = 1; use_wchar = get_mode();
9                 break;
10            case 1:
11                if (use_wchar) {state = 2; read(wbuffer);}
12                else           {state = 3; read( buffer);}
13                return;
14            case 2: handle(wbuffer); return;
15            case 3: handle( buffer); return;
16        } } };
```

Coroutine stack example

```
1 void cr(embo::yield_t<void()>& yield)
2 {
3     bool use_wchar = get_mode();
4     yield();
5     if (use_wchar)
6     {
7         array<wchar_t, 128> wbuffer;
8         read(wbuffer); yield();
9         handle(wbuffer);
10    }
11    else
12    {
13        array<char, 128> buffer;
14        read(buffer); yield();
15        handle(buffer);
16    }
17 }
```

Coroutine stack template example

```
1  template<typename Char>
2  void cr_impl(embo::yield_t<void()> &yield)
3  {
4      array<Char, 128> buffer;
5      read(buffer); yield();
6      handle(buffer);
7  }
8
9  void cr(embo::yield_t<void()>& yield)
10 {
11     bool use_wchar = get_mode();
12     yield();
13     if (use_wchar)
14         cr_impl<wchar_t>(yield);
15     else
16         cr_impl<char>(yield);
17 }
```

- www.github.com/klemens-morgenstern/embo.coroutine

Summary

- www.github.com/klemens-morgenstern/embo.coroutine
- klemens.d.morgenstern@gmail.com

Summary

- www.github.com/klemens-morgenstern/embo.coroutine
- klemens.d.morgenstern@gmail.com
- Any questions?